

Comparison of Virtualization Techniques in Distributed Cloud Environments

Lauri Suomalainen

Master's Thesis
UNIVERSITY OF HELSINKI
Department of Computer Science

Helsinki, October 17, 2016

Tiedekunta — Fakultet — Faculty		Laitos — Institution — Department	
Faculty of Science		Department of Computer Science	
Tekijä — Författare — Author			
Lauri Suomalainen			
Työn nimi — Arbetets titel — Title			
Comparison of Virtualization Techniques in Distributed Cloud Environments			
Oppiaine — Läroämne — Subject			
Computer Science			
Työn laji — Arbetets art — Level	Aika — Datum — Month and year	Sivumäärä — Sidoantal — Number of pages	
Master's Thesis	October 17, 2016	2	
Tiivistelmä — Referat — Abstract			
Abstract			
Avainsanat — Nyckelord — Keywords			
Virtualization, Distributed Systems, Containerization			
Säilytyspaikka — Förvaringsställe — Where deposited			
Muita tietoja — Övriga uppgifter — Additional information			

Contents

1	Introduction	1
2	Background	1
3	System design and implementation	2
4	Experiments and measurements	2
5	Related work	2
6	Future research	2
7	Conclusions	2
	Sources	2

1 Introduction

This is a reference [1]

2 Background

In cloud computing, there are multiple recognized service models which dictate how the users can use the given system and what privileges they are given [4]. In its most limited form, a cloud service is offered to user as a predefined application or a set of applications. The users has some interface for interacting with the applications but is given no control over anything else such as other applications, the operating system the application is running on or network and hardware configurations. This is generally known as Software as a Service or SaaS for short. The most permissible service model is known as IaaS, Infrastructure as a Service. In its archetype the user gets access to all fundamental computing resources, possibly including some network components, and can run arbitrary software including operating systems. The user experience should be similar to that with their personal computers. The user is not allowed to access the underlying cloud infrastructure. Between the two falls Platform as a Service (PaaS). PaaS typically allows user to deploy their own applications along with their dependent libraries, tools, services etc. provided that they are supported by the cloud provider. The user has no control over underlying cloud infrastructure, operating system, storage or network but usually can configure certain settings and possibly choose different supporting services the cloud provider offers. There are also other "aaS" such as Data as a Service (DaaS) and Storage as a Service (SaaS) but they are based on one of the three aforementioned service models or are variations or subsets of them. Sometimes the numerous models are referenced with umbrella terms of XaaS and EaaS meaning Everything as a Service for both or Anything as a Service for the former [3]

Virtualisation in the context of distributed cloud environments usually refers to virtual machines. The core idea is analogous to computer hardware virtualisation. Operating systems offer an interface for the processes to utilize the computer hardware while giving them an illusion that they have all of the hardware for themselves [1]. In reality the resources are shared among many processes. Likewise in cloud environments resources are being share by processes but also by different users running different operating systems, configurations and programs. As with the processes, users are given an impression that they alone have access to the underlying hardware resources, whereas in reality there are multiple users using the same physical machines.

There are several reasons as to why would one prefer a virtualised environment to a non-virtualised one. Obviously in multitenant cloud services it is crucial for the service provider to maximise the use of their hardware

resources. Thus it is imperative for the provider to try to share the limited hardware resources among as many users' virtualised environments as possible. Otherwise every user would need their own physical machine in the system which would both require more resources per user and leave resources underused. From the fault tolerance perspective, using virtual machines in a distributed environment decreases their dependency on the underlying physical hardware [2]. That is because in virtual machine architectures which support live migration operating system instances can be seamlessly moved from one physical machine to another. This also helps the load-balancing in the distributed system and allows low-level and physical maintenance of the hardware without considerably disrupting the usage of the system.

3 System design and implementation

4 Experiments and measurements

5 Related work

6 Future research

7 Conclusions

Sources

- [1] Arpaci-Dusseau, Remzi H. and Arpaci-Dusseau, Andrea C.: *Operating Systems: Three Easy Pieces*. Arpaci-Dusseau Books, 0.91 edition, May 2015.
- [2] Clark, Christopher, Fraser, Keir, H, Steven, Hansen, Jacob Gorm, Jul, Eric, Limpach, Christian, Pratt, Ian, and Warfield, Andrew: *Live migration of virtual machines*. In *In Proceedings of the 2nd ACM/USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, pages 273–286, 2005.
- [3] Duan, Y., Fu, G., Zhou, N., Sun, X., Narendra, N. C., and Hu, B.: *Everything as a service (xaas) on the cloud: Origins, current and future trends*. In *2015 IEEE 8th International Conference on Cloud Computing*, pages 621–628, June 2015.
- [4] Mell, Peter M. and Grance, Timothy: *Sp 800-145. the nist definition of cloud computing*. Technical report, Gaithersburg, MD, United States, 2011.