

Comparison of Virtualization Techniques in Distributed Cloud Environments

Lauri Suomalainen

Master's Thesis
UNIVERSITY OF HELSINKI
Department of Computer Science

Helsinki, November 7, 2016

Tiedekunta — Fakultet — Faculty		Laitos — Institution — Department	
Faculty of Science		Department of Computer Science	
Tekijä — Författare — Author			
Lauri Suomalainen			
Työn nimi — Arbetets titel — Title			
Comparison of Virtualization Techniques in Distributed Cloud Environments			
Oppiaine — Läroämne — Subject			
Computer Science			
Työn laji — Arbetets art — Level	Aika — Datum — Month and year	Sivumäärä — Sidoantal — Number of pages	
Master's Thesis	November 7, 2016	7	
Tiivistelmä — Referat — Abstract			
Abstract			
Avainsanat — Nyckelord — Keywords			
Virtualization, Distributed Systems, Containerization			
Säilytyspaikka — Förvaringsställe — Where deposited			
Muita tietoja — Övriga uppgifter — Additional information			

Contents

1	Introduction	1
2	Background	1
2.1	Virtualisation Techniques	2
2.1.1	Full virtualisation	2
2.1.2	Hardware-Layer virtualisation	3
2.1.3	Container-based virtualisation	4
2.1.4	Paravirtualisation	5
3	System design and implementation	6
4	Experiments and measurements	6
5	Related work	6
6	Future research	6
7	Conclusions	6
	Sources	6

1 Introduction

This is a reference [7]

2 Background

In cloud computing, there are multiple recognized service models which dictate how the users can use the given system and what privileges they are given [11]. In its most limited form, a cloud service is offered to user as a predefined application or a set of applications. The users has some interface for interacting with the applications but is given no control over anything else such as other applications, the operating system the application is running on or network and hardware configurations. This is generally known as Software as a Service or SaaS for short. The most permissible service model is known as IaaS, Infrastructure as a Service. In its archetype the user gets access to all fundamental computing resources, possibly including some network components, and can run arbitrary software including operating systems. The user experience should be similar to that with their personal computers. The user is not allowed to access the underlying cloud infrastructure. Between the two falls Platform as a Service (PaaS). PaaS typically allows user to deploy their own applications along with their dependent libraries, tools, services etc. provided that they are supported by the cloud provider. The user has no control over underlying cloud infrastructure, operating system, storage or network but usually can configure certain settings and possibly choose different supporting services the cloud provider offers. There are also other "aaS" such as Data as a Service (DaaS) and Storage as a Service (SaaS) but they are based on one of the three aforementioned service models or are variations or subsets of them. Sometimes the numerous models are referenced with umbrella terms of XaaS and EaaS meaning Everything as a Service for both or Anything as a Service for the former [9].

Virtualisation in the context of distributed cloud environments usually refers to virtual machines. The core idea is analogous to computer hardware virtualisation. Operating systems offer an interface for the processes to utilize the computer hardware while giving them an illusion that they have all of the hardware for themselves [7]. In reality the resources are shared among many processes. Likewise in cloud environments resources are being share by processes but also by different users running different operating systems, configurations and programs. As with the processes, users are given an impression that they alone have access to the underlying hardware resources, whereas in reality there are multiple users using the same physical machines.

There are several reasons as to why would one prefer a virtualised environment to a non-virtualised one. Obviously in multitenant cloud services it is crucial for the service provider to maximise the use of their hardware

resources. Thus it is imperative for the provider to try to share the limited hardware resources among as many users' virtualised environments as possible. Otherwise every user would need their own physical machine in the system which would both require more resources per user and leave resources underused. From the fault tolerance perspective, using virtual machines in a distributed environment decreases their dependency on the underlying physical hardware [8]. That is because in virtual machine architectures which support live migration of operating system instances can be seamlessly moved from one physical machine to another. This also helps the load-balancing in the distributed system and allows low-level and physical maintenance of the hardware without considerably disrupting the usage of the system. A end-user also has many reasons to use virtualised cloud services. User only needs a lightweight computer with an internet connection to perform computationally challenging tasks in the cloud back-end. Similarly devices with little storage capacity can leverage from a cloud service's vast storage space. Some users would like to use applications and programs not native to their operating system of choice making another virtualised OS a convenient option [7]. Virtualised environment allows software developers to test and debug their software with many different settings, as virtualised environments can have different operating systems and available hardware resources. Naturally this also allows emulating completely different devices [10].

2.1 Virtualisation Techniques

Traditionally virtualisation has referred to a software abstraction layer residing between the computer hardware and the operating system. [12] This layer has been called Virtual Machine Monitor (VMM) or more recently a hypervisor and it hides and abstracts the computing resources from the OS, allowing multiple OSs to run simultaneously on the same hardware. There are multiple ways to run hypervisor-based virtualisation. Lately a technology called container-based virtualisation has been gaining popularity. Instead of emulating whole hardware, containers make use of features provided by the host operating system to isolate processes from each other and other containers [10].

2.1.1 Full virtualisation

In full virtualisation, the hypervisor runs on top of the host OS. The guest OSs run on top of the hypervisor which in turn emulates the underlying real hardware to them. The hypervisors running on top of the host OS are generally referred as *Type 2 Hypervisors* [10]. The guest OSs can be arbitrary. Figure 1 shows the full virtualisation architecture with the hypervisor running on top of the Host OS and Guest OSs on top of the hypervisor using their emulated hardware.

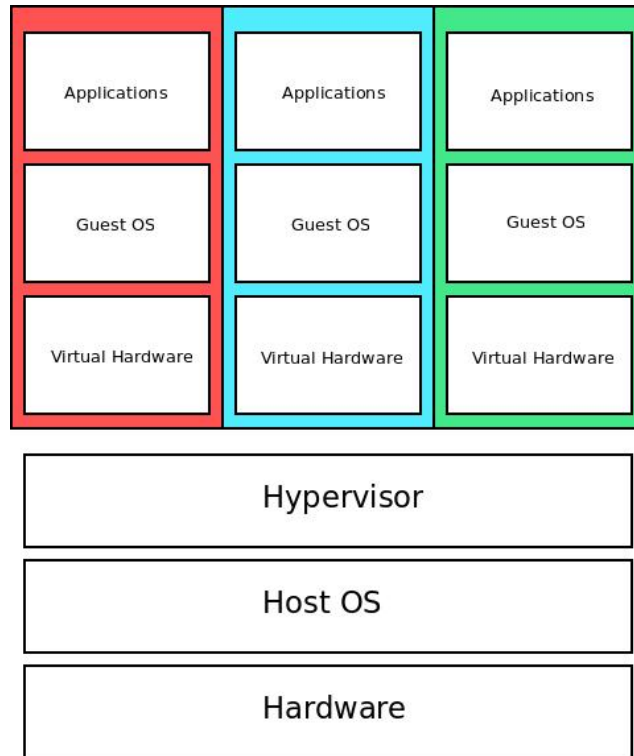


Figure 1: Full virtualisation architecture

The main advantage of full virtualisation is that it is easy to deploy and should not pose problems to an average user but the virtualisation overhead results in significantly reduced performance when compared to running directly on hardware. Popular examples of full virtualisation applications are Oracle’s *VirtualBox*[3] and *VMware Workstation*[5].

2.1.2 Hardware-Layer virtualisation

Hardware-Layer virtualisation is also a type of full virtualisation, but unlike Type 2 hypervisors, the so called *Type 1 Hypervisors* (also *native* and *bare metal*) run directly on hardware. As seen on figure 2 there’s no Host OS per se. Instead the Guest OSs access to hardware resources is controlled by the hypervisor.

Running directly on hardware, Hardware-Layer virtualisation techniques suffer less performance overhead than their OS-layer counterparts. On the other hand, Type 2 hypervisors being essentially applications themselves can be ran in parallel on the host OS whereas Type 1 hypervisors can not. For an average user, setting up a Type 1 hypervisor can be more difficult than Type 2. Commercial examples of Type 1 Hypervisors include Microsoft’s *Hyper-V*[2] and VMware’s *VSphere* [4].

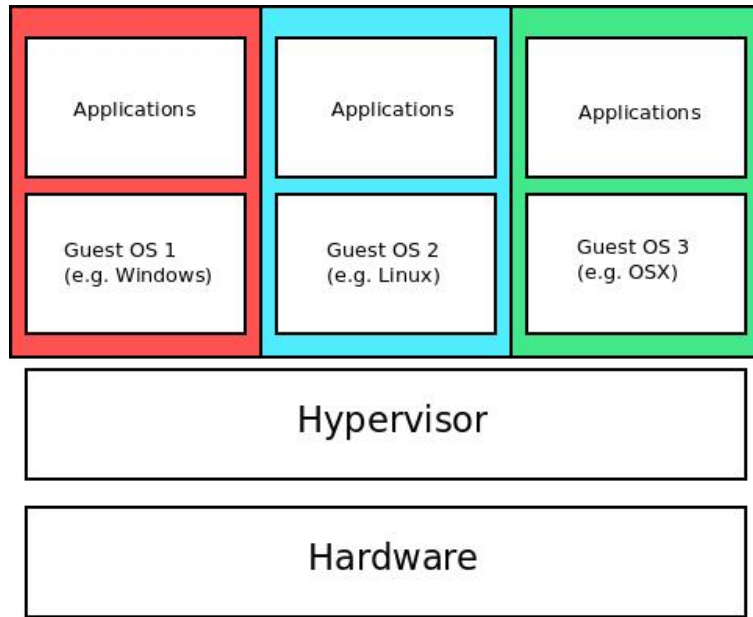


Figure 2: Hardware-Layer virtualisation architecture

2.1.3 Container-based virtualisation

Instead of virtualising the underlying hardware, container-based virtualisation also known as OS-Layer virtualisation [12] focuses on user space and allows running multiple operating systems in parallel as applications using the same kernel as the host operating system. A prime example of a popular container-based virtualisation platform is *Docker* [1] which leverages on native Linux kernel features to virtualise and isolate OS instances. Figure 3 shows a container-based virtualisation architecture in which containerised environments are running operating systems on host OS's kernel.

Container-based virtualisation does not need to emulate the hardware as containers communicate directly with the host kernel [10] and are thus very fast to start. They also do not require all of the components a fully virtualised environment would need to run and therefore their resource footprint is minimal when compared to hypervisor-based virtualisation techniques. The obvious drawback of the technique is that the kernel of the virtualised OS has to be the same as that of Host OS e.g. In a situation depicted in figure 3 operating systems based on Linux kernel could be ran on Ubuntu Host OS but OSs like Windows or OSX could not. On certain virtualisation platforms resource-intensive containers can also affect other containers detrimentally as the shared host OS's kernel is forced to spend its execution time on handling the instructions from the stressed container [13].




Figure 3: Container-based architecture

2.1.4 Paravirtualisation

Paravirtualisation differs from full virtualisation by requiring the Guest OS to be modified in order to accommodate the virtual environment in which it is ran. Otherwise the architecture is similar to that of full virtualisation, but with thinner hypervisor allowing performance close to that of a non-virtualised environment. A well-known example of a paravirtualisation hypervisor is *Xen*[6].

3 System design and implementation

4 Experiments and measurements

5 Related work

6 Future research

7 Conclusions

Sources

- [1] *Docker*. <https://www.docker.com>. Accessed: 3.11.2016.
- [2] *Microsoft hyper-v*. [https://technet.microsoft.com/en-us/library/mt169373\(v=ws.11\).aspx](https://technet.microsoft.com/en-us/library/mt169373(v=ws.11).aspx). Accessed: 7.11.2016.
- [3] *Oracle virtualbox*. <https://www.virtualbox.org/>. Accessed: 1.11.2016.
- [4] *Vmware vsphere hypervisor*. <http://www.vmware.com/products/vsphere-hypervisor.html>. Accessed: 7.11.2016.
- [5] *VMware workstation*. <http://www.vmware.com/products/workstation.html>. Accessed: 1.11.2016.
- [6] *Xen project*. <https://www.xenproject.org/1>. Accessed: 7.11.2016.
- [7] Arpaci-Dusseau, Remzi H. and Arpaci-Dusseau, Andrea C.: *Operating Systems: Three Easy Pieces*. Arpaci-Dusseau Books, 0.91 edition, May 2015.
- [8] Clark, Christopher, Fraser, Keir, H, Steven, Hansen, Jacob Gorm, Jul, Eric, Limpach, Christian, Pratt, Ian, and Warfield, Andrew: *Live migration of virtual machines*. In *In Proceedings of the 2nd ACM/USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, pages 273–286, 2005.
- [9] Duan, Y., Fu, G., Zhou, N., Sun, X., Narendra, N. C., and Hu, B.: *Everything as a service (xaas) on the cloud: Origins, current and future trends*. In *2015 IEEE 8th International Conference on Cloud Computing*, pages 621–628, June 2015.
- [10] Eder, Michael: *Hypervisor-vs. container-based virtualization*. Future Internet (FI) and Innovative Internet Technologies and Mobile Communications (IITM), 1, 2016.

- [11] Mell, Peter M. and Grance, Timothy: *Sp 800-145. the nist definition of cloud computing*. Technical report, Gaithersburg, MD, United States, 2011.
- [12] Sahoo, J., Mohapatra, S., and Lath, R.: *Virtualization: A survey on concepts, taxonomy and associated security issues*. In *Computer and Network Technology (ICCNT), 2010 Second International Conference on*, pages 222–226, April 2010.
- [13] Xavier, M. G., Neves, M. V., Rossi, F. D., Ferreto, T. C., Lange, T., and Rose, C. A. F. De: *Performance evaluation of container-based virtualization for high performance computing environments*. In *2013 21st Euromicro International Conference on Parallel, Distributed, and Network-Based Processing*, pages 233–240, Feb 2013.